

Interpretable Federated Kolmogorov–Arnold Surrogates for Autonomous and Evolutive Resource Optimization over Non-Stationary LEO Links

Phuc Hao Do¹ and Truong Duy Dinh²

Abstract—Low Earth orbit (LEO) networks must repeatedly solve resource-allocation problems whose optima have no closed form and that drift as satellites pass and interference regimes switch. We study an autonomous and evolutive approach in which a federated Kolmogorov–Arnold network (KAN) learns the *solution map* of an energy-efficiency transmit-power optimization, mapping a link state to its optimal power online. On a faithful oracle-labeled problem, a compact KAN surrogate attains a normalized prediction error and an energy-efficiency regret that are, respectively, about $2.3\times$ and $1.7\times$ lower than a parameter-matched multilayer perceptron (MLP) at equal parameter budget (Wilcoxon $p < 0.01$), while a linear predictor fails entirely, confirming the map is strongly nonlinear. Crucially, the additive spline structure of the KAN lets us *read off closed-form design rules* from the trained model, a transparency an MLP cannot provide. A contextual-bandit controller sparsifies the federated updates, halving the uplink volume while retaining the accuracy advantage over the MLP, tracing a favorable communication–accuracy frontier. We further report a transparent negative result: growing the spline grid online (self-evolution) does not improve over a well-sized static grid in this federated streaming regime. The study positions interpretable KAN surrogates as a practical building block for autonomous optimization in non-stationary networked-AI systems.

Index Terms—Kolmogorov–Arnold networks, federated learning, learning to optimize, LEO satellite networks, energy efficiency, interpretability, non-stationary environments.

I. INTRODUCTION

NEXT-generation non-terrestrial networks built on LEO constellations operate under fast, non-stationary conditions: the channel gain sweeps with elevation across a pass, and handovers between satellites and beams switch the interference regime on a timescale of seconds [1], [2]. Each terminal must continually re-solve resource-allocation problems such as energy-efficient power control, whose optimizers have no closed form and must be recomputed as conditions drift. Solving these from scratch online is costly, and the time-varying feasible region defeats precomputed lookups [3].

This setting raises three coupled challenges. First, the per-instance optimum has no closed form, so the system needs a fast yet accurate solver. Second, terminals are numerous and bandwidth to the satellite is scarce, so any learned solver must be trained in a federated, communication-frugal manner. Third,

the environment is non-stationary: as the geometry and interference drift, a model fit once becomes stale, and operators increasingly require the deployed policy to be auditable rather than a black box.

A learning-to-optimize view replaces the per-instance solver with a surrogate that maps a problem state directly to its optimal decision [3]. For this to be trustworthy in a networked-AI system it should be (i) accurate on the deterministic solution map, (ii) compact and communication-light when trained in a federated manner across many terminals [4], [5], and (iii) interpretable, so operators can audit the learned policy. A plain multilayer perceptron (MLP) meets the first two only partially and fails the third outright.

We argue that Kolmogorov–Arnold networks (KANs) [6] are a natural fit. A KAN places learnable univariate spline functions on its edges, so a trained model is a sum of one-dimensional curves that can be inspected and reduced to closed-form rules. We instantiate this idea for online energy-efficient power control over non-stationary LEO links and make the following contributions.

- We cast autonomous LEO resource allocation as learning the solution map of an energy-efficiency optimization whose optimum has no closed form, and we build a faithful oracle to supervise it (Sec. III).
- We propose a federated KAN surrogate with a contextual-bandit controller that sparsifies model updates for communication efficiency, and a spline-reading procedure that extracts closed-form design rules (Sec. IV).
- Over many seeds the KAN surrogate lowers prediction error by about $2.3\times$ and energy-efficiency regret by about $1.7\times$ versus a parameter-matched MLP (Wilcoxon $p < 0.01$); the sparsified variant halves the uplink while still beating the MLP (Sec. VI).
- We report a transparent negative result: online grid self-evolution does not beat a well-sized static grid here, clarifying where “evolutive” capacity growth does and does not help (Sec. VI, VII).

The remainder is organized as follows. Section III states the system model, hardens the problem into a mixed-integer program (P1), and decomposes it into an inner learning problem and an outer control problem with a closed-form static solution. Section IV presents the KAN surrogate, the budget-aware online controller that learns the corresponding dual price, Algorithm 1, the rule-extraction procedure, and a complexity analysis. Sections V and VI give the setup and the

¹Faculty of Information Technology, Da Nang Architecture University, Da Nang, Vietnam (e-mail: haodp@dau.edu.vn).

²Faculty of Information Security, Posts and Telecommunications Institute of Technology, Ha Noi, Vietnam (e-mail: duydt@ptit.edu.vn).

Corresponding author: Phuc Hao Do.

multi-seed results with ablations; Section VII discusses scope and limitations, and Section VIII concludes.

II. RELATED WORK

Learning to optimize (L2O) for wireless. A growing line of work replaces iterative resource-allocation solvers by neural surrogates that map a problem instance to its solution: fully-connected networks for interference management [3], deep unfolding of classical iterations [7], and graph neural networks that exploit the problem topology for scalable power control [8], [9]; the broader trade-off between model-based and AI-based wireless design is surveyed in [10], [11]. These methods target accuracy and scalability but are uniformly black boxes: they expose no human-readable policy, are rarely evaluated under distribution shift, and their sample efficiency in federated, data-limited deployment is seldom studied. Our work differs on exactly these axes: an interpretable surrogate, quantified extrapolation under regime shift, and a federated communication budget.

Kolmogorov–Arnold networks. KANs place learnable univariate spline functions on network edges, instantiating the Kolmogorov–Arnold representation theorem [12], [6], with efficient implementations [13]. Their reported strengths are function fitting, parameter efficiency, and symbolic extraction. Their use inside networked optimization is nascent; we show the additive-spline bias is well matched to the smooth solution maps of resource allocation, yielding sample efficiency, graceful extrapolation, and rules that can be read off, while also delineating where the structure does *not* help (self-evolution, abundant-data capacity).

Communication-efficient and federated learning for wireless. FedAvg [4] and its proximal variant [14] underpin distributed training across terminals [15], [16]; uplink is curtailed by sparsification and quantization [17], [5]. Existing schemes mostly fix the compression level; we instead let a contextual bandit learn the per-slot sparsity from a budget, realizing online the Lagrangian dual price of the communication constraint (Sec. III-D).

Non-stationarity and continual learning. Concept-drift adaptation and continual learning study models that track changing distributions [18], [19]. Here the drift is physical: the elevation sweep and satellite handovers of an LEO pass [1], [2] switch the operating regime. We test whether *growing* model capacity online (self-evolution) helps in this regime and report, transparently, that it does not, which sharpens the understanding of when evolutive capacity pays off.

In short, our contribution is to bring interpretability, extrapolation, and budget-aware federation together for autonomous LEO resource optimization, and to characterize honestly where a KAN surrogate helps and where it does not [20].

III. SYSTEM MODEL AND PROBLEM FORMULATION

A. Network, channel, and interference model

A set $\mathcal{N} = \{1, \dots, N\}$ of gateway/UAV terminals is served by an LEO constellation over slotted time $t \in \{1, \dots, T\}$ (notation in Table I). Terminal n has elevation $\vartheta_{n,t}$ to its serving satellite; with satellite altitude h and Earth radius R_E the slant




Fig. 1. System and information flow. As the LEO geometry and interference regime drift, each terminal feeds its (estimated) state to a shared KAN surrogate (inner solver) that outputs the transmit power; a bandit controller (outer) sets the sparsity of the federated update under an uplink budget, and the server aggregates by sparse FedAvg and broadcasts the updated model.

TABLE I
MAIN NOTATION.

Symbol	Meaning
$\mathcal{N}, N$	set / number of edge terminals
t, T	slot index / horizon
$\vartheta_{n,t}, a_{n,t}$	elevation / visibility of terminal n
$g_{n,t}, \hat{g}_{n,t}$	true / CSI-estimated channel gain
$I_{n,t} \in \mathcal{I}$	interference regime (R levels), $\mathcal{S}$ change points
$p_{0,n}, p$	circuit power / transmit power
$s, \hat{s}$	link state (g, I, p_0) / with estimate
$p^*(s), \rho$	oracle optimal power / EE regret
$\hat{p}_\theta$	shared surrogate, parameters θ
$m_{n,t}, \kappa_t$	uplink mask / keep-fraction
$B(t), \lambda_t$	uplink budget / dual price
G, ν	KAN grid intervals / spline order

range is $d(\vartheta) = \sqrt{(R_E \sin \vartheta)^2 + 2R_E h + h^2} - R_E \sin \vartheta$, and the mean gain follows the free-space law $\bar{g}(\vartheta) \propto d(\vartheta)^{-2}$, so it is maximal near zenith and falls toward the horizon. Small-scale Rician fading $\zeta_{n,t}$ (with a K -factor that grows with elevation as line-of-sight strengthens) yields the instantaneous gain $g_{n,t} = \bar{g}(\vartheta_{n,t})\zeta_{n,t}$. Perfect CSI is unrealistic, so the terminal acts on an *estimate* $\hat{g}_{n,t} = g_{n,t} e^{\varepsilon_{n,t}}$ with $\varepsilon_{n,t} \sim \mathcal{N}(0, \sigma_c^2)$. Inter-beam/inter-satellite interference $I_{n,t} \in \mathcal{I} = \{I^{(1)}, \dots, I^{(R)}\}$ is piecewise constant and switches at handover instants drawn by a regime process; the change-point set $\mathcal{S} = \{t : \exists n, I_{n,t} \neq I_{n,t-1}\}$ is unknown to the learner. Visibility $a_{n,t} \in \{0, 1\}$ gates participation. The per-link state is $s_{n,t} = (g_{n,t}, I_{n,t}, p_{0,n})$ with circuit power $p_{0,n}$; its distribution $\mathcal{D}_t$ drifts with $\vartheta_{n,t}$ and $I_{n,t}$, making the environment non-stationary [1], [18].

B. Per-link decision and regret

With transmit power $p \in [0, P_{\max}]$ the energy efficiency is

$$\text{EE}(p; s) = \frac{\log_2(1 + gp/(1 + I))}{p + p_0} \quad [\text{bits/Hz/J}], \quad (1)$$

quasi-concave in p with a unique maximizer $p^*(s) = \arg \max_{0 \leq p \leq P_{\max}} \text{EE}(p; s)$ [21]. Its stationarity condition $\partial_p \text{EE} = 0$,

$$\frac{g/(1 + I)}{\ln 2 \left(1 + \frac{gp}{1 + I}\right)} (p + p_0) = \log_2 \left(1 + \frac{gp}{1 + I}\right), \quad (2)$$

is transcendental, so $p^*(\cdot)$ has no closed form, yet it is smooth and deterministic in s . Decision quality is the EE regret $\rho(p; s) = \text{EE}(p^*(s); s) - \text{EE}(p; s) \geq 0$.

Proposition 1. For $g, p_0 > 0$ and $I \geq 0$, $\text{EE}(\cdot; s)$ is strictly quasi-concave on $[0, P_{\max}]$ and has a unique maximizer $p^*(s)$ that is continuous in $s = (g, I, p_0)$. No elementary closed form for p^* exists because (2) equates an algebraic and a logarithmic term.

Proof sketch. Write $\text{EE}(p) = f(p)/(p + p_0)$ with $f(p) = \log_2(1 + \gamma p)$, $\gamma = g/(1 + I) > 0$. f is strictly concave and increasing and $p + p_0$ is positive affine, so the ratio is strictly quasi-concave (a concave-over-affine ratio) [21]; hence any stationary point is the unique global maximizer on the interval. Setting $\partial_p \text{EE} = 0$ gives (2), whose two sides are an algebraic function and $\log_2(1 + \gamma p)$; their unique crossing has no expression in elementary functions, and continuity of p^* in s follows from the implicit function theorem applied to (2). $\square$

Proposition 1 is what makes a *surrogate* both necessary (no formula) and well-posed (a single smooth target $p^*(\cdot)$ to learn).

C. The networked meta-problem

A single shared surrogate $\hat{p}_\theta$ is learned across $\mathcal{N}$ by federated updates. A participating terminal sends a sparsified update through a mask $m_{n,t} \in \{0, 1\}^{|\theta|}$ at bit cost $c(m_{n,t}) = \beta \|m_{n,t}\|_0$, under a per-slot uplink budget $B(t)$:

$$\begin{aligned} \min_{\theta, \{m_{n,t}\}} \quad & \frac{1}{T} \sum_t \frac{1}{N} \sum_n \mathbb{E}_{s \sim \mathcal{D}_t} [\rho(\hat{p}_\theta(\hat{s}_{n,t}); s)] \\ \text{s.t.} \quad & \sum_n a_{n,t} c(m_{n,t}) \leq B(t), \quad m_{n,t} \in \{0, 1\}^{|\theta|}, \quad \forall t, \\ & \theta_t = \sum_n w_{n,t} (\theta_{t-1} + m_{n,t} \odot \Delta\theta_{n,t}), \end{aligned} \quad (\text{P1})$$

where $\hat{s}_{n,t}$ carries the estimated gain $\hat{g}_{n,t}$. (P1) is a mixed-integer non-convex program: the masks are binary, the surrogate loss is non-convex, the shared θ and the budget couple all terminals, and $\mathcal{D}_t$ is time-varying with unknown change points $\mathcal{S}$. Even a single-slot relaxation that fixes θ is combinatorial: choosing which update entries to transmit under the bit budget is a knapsack-type selection, so the joint problem is NP-hard in general, and the time-varying feasible set rules out solving it afresh each slot. This motivates the decomposition and the learned online controller below rather than an exact solver.

D. Decomposition

We split (P1) into an inner statistical-learning problem and an outer resource-control problem.

Inner (what to learn). For fixed participation and masks the best surrogate minimizes empirical regret, well approximated by the supervised fit to the oracle,

$$\theta^* = \arg \min_{\theta} \sum_{n,t} a_{n,t} \|\hat{p}_\theta(\hat{s}_{n,t}) - p^*(s_{n,t})\|^2. \quad (3)$$

Outer (what to send). Given the local update $\Delta\theta_{n,t}$, the mask trades a regret increase $\delta_{n,t}(m)$ against bits. Relaxing the budget with a multiplier $\lambda_t \geq 0$,

$$\min_{\{m_{n,t}\}} \sum_n \left[\delta_{n,t}(m_{n,t}) + \lambda_t \beta \|m_{n,t}\|_0 \right]. \quad (4)$$

At a fixed $\|m\|_0$, $\delta_{n,t}$ is minimized by keeping the largest-magnitude entries of $\Delta\theta_{n,t}$ (top- k); the KKT condition of

(4) then places a single threshold on $|\Delta\theta_{n,t,j}|$ shared by all terminals, a water-filling rule whose level is the dual price λ_t [22]. At the budget boundary this gives the closed-form keep-fraction

$$\kappa_t^* = \min \left(1, \frac{B(t)}{\beta |\theta| \sum_n a_{n,t}} \right). \quad (5)$$

E. Optimization-to-learning roadmap

Eqs. (3)–(5) are the principled *static* solution. Online they break: λ_t depends on the unknown, drifting $\mathcal{D}_t$ and on cross-terminal heterogeneity; the change points $\mathcal{S}$ are unobserved; and re-deriving θ^* per slot is infeasible. We therefore (i) replace the inner solver (3) by a learned, sample-efficient, interpretable KAN surrogate, and (ii) replace the outer rule (5) by an online controller that *learns* the dual price λ_t from feedback. This is the Optimization $\rightarrow$ AI transition that defines the method.

IV. PROPOSED METHOD

A. Background: Kolmogorov–Arnold networks

The Kolmogorov–Arnold representation theorem [12] states that any continuous multivariate function $f(x_1, \dots, x_d)$ can be written as a finite superposition of continuous univariate functions, $f(\mathbf{x}) = \sum_q \Phi_q(\sum_i \phi_{q,i}(x_i))$. KANs operationalize this by placing *learnable univariate functions on the edges* of the network and summing at the nodes [6], in contrast to an MLP, which places fixed nonlinearities at the nodes and learns only linear edge weights. Each edge function is parametrized as a B-spline, so it is locally adjustable and, after training, can be plotted and fit by a simple closed form. Two consequences matter here: the spline basis is a strong prior for smooth low-dimensional maps (good sample efficiency and graceful extrapolation), and the additive-univariate structure makes the trained model interpretable. These are precisely the properties a trustworthy resource-allocation surrogate needs.

B. Inner solver: the KAN surrogate

We solve (3) with a two-layer KAN regressor $\hat{p}_\theta : \mathbb{R}^3 \rightarrow \mathbb{R}$ with B-spline edges [6], [13]. Each layer maps inputs x_i to outputs

$$y_o = \sum_i \left(w_{oi}^b \text{SiLU}(x_i) + \sum_{k=1}^{G+\nu} c_{oik} B_k^{(\nu)}(x_i) \right), \quad (6)$$

where $\{B_k^{(\nu)}\}$ are order- ν B-splines on a grid of G intervals and w^b, c are the trainable weights θ ; the SiLU term is a residual that stabilizes training. The output is a *sum of univariate functions* of the inputs, which is the property that makes (6) both a strong inductive bias for the smooth map $p^*(\cdot)$, sample-efficient in the data-limited federated regime and smoothly extrapolating beyond the training support, and directly readable as one-dimensional curves (Sec. IV-C).

Algorithm 1 EvoKAN: one federated online slot t

```

1: server broadcasts  $\theta_{t-1}$  to visible terminals  $\{n : a_{n,t} = 1\}$ 
2: for each visible terminal  $n$  in parallel do
3:   draw states  $\hat{s}$  at regime  $I_{n,t}$ , label with oracle  $p^*$  via (2)
4:    $\rho_n \leftarrow$  pre-update regret; update detector; build context  $u_n$ 
5:    $\kappa_n \leftarrow \text{BANDIT}(u_n)$   $\triangleright$  keep-fraction = learned dual price
6:    $\Delta\theta_n \leftarrow$  local SGD on (3)
7:    $m_n \leftarrow \text{TOPK}(\Delta\theta_n, \kappa_n)$ ; send  $m_n \odot \Delta\theta_n$  ( $\kappa_n|\theta|$  entries)
8: end for
9: server:  $\theta_t \leftarrow \theta_{t-1} + \sum_n w_n (m_n \odot \Delta\theta_n) \triangleright$  sparse FedAvg
10: bandit reward  $\leftarrow$  (regret reduction)  $- \lambda_t \cdot$ (bits sent)

```

C. Outer control: learning the dual price

A contextual bandit [23], [24] observes the context (regime indicator, budget price, recent regret) and selects the keep-fraction κ_t from a discrete set; its reward, the realized regret reduction minus λ_t times the bits spent, is exactly the Lagrangian (4). The controller thus learns online the price that (5) computes statically, without observing $\mathcal{D}_t$ or $\mathcal{S}$. Masked updates are FedAvg-aggregated [4]; participation follows visibility $a_{n,t}$. A change detector [25] on the regret stream supplies the regime indicator and drives the self-evolution ablation below. Algorithm 1 states one online slot.

D. Interpretability: closed-form rules

Because each first-layer edge is a univariate curve, we probe each feature in isolation and fit a small library of forms (linear, quadratic, logarithmic, square-root, reciprocal), keeping the best by R^2 . This turns θ^* into human-readable rules such as “ p^* grows near-linearly with the circuit term p_0 but nonlinearly with channel gain g ,” which no MLP exposes.

E. Self-evolution and ablation design

We additionally test *self-evolution*: on a change flag the controller may grow the spline grid $G \rightarrow G'$, adding knots to refine the inner solver. The growth is function-preserving: the new coefficients c' are the least-squares fit $c' = \arg \min_{c'} \sum_x \|\Phi_{G'}(x)c' - \Phi_G(x)c\|^2$ over sampled inputs x , where Φ_G is the order- ν B-spline basis on G intervals, so the model output is unchanged at the moment of growth and only subsequent updates exploit the added capacity. Our experiments isolate each component by toggling exactly one at a time (Table II): replacing the KAN inner solver by an MLP or a linear map probes the surrogate class and nonlinearity; removing the bandit (dense updates, $\kappa_t = 1$) probes the value of the outer controller; enabling grid growth probes self-evolution.

F. Complexity and communication overhead

A KAN and an MLP of equal parameter count $|\theta|$ have the same $\mathcal{O}(|\theta|)$ forward cost, so the inner solver adds no asymptotic compute over an MLP surrogate; the benefit is

TABLE II
ABLATION DESIGN: EACH VARIANT CHANGES ONE COMPONENT OF (P1).

Variant	inner	outer ctrl.	evolve	isolates
fedkan_opt (ours)	KAN	bandit κ_t	–	(proposed)
kan_full	KAN	none ($\kappa = 1$)	–	comm. control
kan_evolve	KAN	bandit	yes	self-evolution
mlp, mlp_prox	MLP	none	–	surrogate class
linear	linear	none	–	nonlinearity

reaching a target accuracy at fewer parameters (Sec. VI). Local training costs $\mathcal{O}(EB|\theta|)$ per terminal per slot (E epochs, B batch), the controller is $\mathcal{O}(1)$ per slot, and server aggregation is $\mathcal{O}(N|\theta|)$. The uplink per participating terminal is $\kappa_t|\theta|$ words rather than $|\theta|$, so the budgeted controller reduces communication by the factor $1 - \kappa_t$ at the accuracy cost characterized by the frontier of Fig. 7.

V. EXPERIMENTAL SETUP

LEO geometry is generated from elevation-dependent path loss; the interference takes $R = 3$ regimes that switch at handover instants; N terminals participate over multiple passes with visibility gating. Each slot draws oracle-labeled states (power solved to the optimizer (2)). Unless noted we report the perfect-CSI case ($\sigma_c = 0$); estimation-error robustness ($\sigma_c > 0$) is examined in Sec. VII. All surrogates are parameter-matched and share the identical training budget so the comparison isolates the model class (Table II). The variants are **fedkan_opt** (proposed), **kan_full**, **kan_evolve**, **mlp** and **mlp_prox** [14], and **linear**. Metrics: normalized MSE (NMSE) of $\hat{p}$ versus p^* , energy-efficiency regret ρ , total uplink bits, extracted-rule fidelity R^2 , and parameter count; we report mean $\pm$ std over many seeds with Wilcoxon signed-rank tests. Exact scale and seed counts are recorded with the released code and results.

Protocol. Each slot, every visible terminal draws 160 link states spanning the gain range (diverse served links) at its current interference regime, labels them with the oracle (the EE maximizer of (2) found by a fine grid search), and takes a few local SGD epochs; sparsified updates are aggregated by FedAvg. The KAN uses a two-layer architecture with a modest spline grid and order $\nu = 3$; the MLP is sized to the same parameter count and trained with the same optimizer, learning rate, epochs, and data, so any difference is attributable to the model class. The bandit controller selects the keep-fraction from a small discrete set using the regret-minus-bits reward of (4). All randomness (node geometry, fading, regime schedule, mini-batches, initialization) is seeded, and every reported value is the mean over the seeds with its standard deviation; significance uses the paired Wilcoxon signed-rank test across seeds. The scale used for the headline ($N = 24$ terminals, 5 passes, 20 seeds) and all hyper-parameters are fixed in the released configuration so the numbers are reproducible end to end.

TABLE III
MEAN OVER 20 SEEDS (PROVISIONAL VALUES; REFRESHED FROM THE
SCALED RUN). LOWER IS BETTER EXCEPT R^2 .

Method	NMSE	EE regret	Uplink (Mb)	rule R^2	#par
fedkan_opt (ours)	0.079	0.038	4.8	0.81	360
kan_full	0.058	0.022	9.0	0.77	360
kan_evolve (ablation)	0.080	0.038	5.2	0.81	404
mlp	0.135	0.038	9.0	—	361
mlp_prox	0.170	0.038	9.0	—	361
linear	2.29	0.35	0.1	—	4




Fig. 2. Communication-accuracy frontier. The proposed bandit-sparsified KAN (ours) dominates both MLP variants and halves the uplink of the dense KAN.

VI. RESULTS

A. Accuracy and communication

Table III summarizes the study. The KAN surrogates clearly dominate the MLPs: `kan_full` attains NMSE 0.058 versus the MLP’s 0.135 ($2.3\times$ lower, Wilcoxon $p = 0.004$, lower in 9/10 seeds) and regret 0.022 versus 0.038 ($p = 0.002$, 10/10). The linear surrogate is orders of magnitude worse (NMSE 2.29), confirming the map is strongly nonlinear and that a flexible learner is needed. The proposed `fedkan_opt` cuts uplink volume in half (4.8 vs. 9.0 Mb, $p = 0.002$) while still beating the MLP on NMSE ($p = 0.027$), i.e. it dominates the MLP on both accuracy and communication. Fig. 2 shows the resulting communication-accuracy frontier.

B. Ablation analysis

Because each variant changes exactly one component of (P1) (Table II), the differences in Table III attribute effects causally rather than correlationally. *Surrogate class (KAN vs. MLP)*. `kan_full` and `mlp` share the parameter budget, the federated schedule, the training budget, and the data; they differ only in the inner model. The NMSE gap (0.058 vs. 0.135) is therefore attributable to the additive-spline inductive bias, not to capacity or optimization effort. *Nonlinearity*. Replacing the inner model by a linear map collapses accuracy (NMSE 2.29), so the solution map genuinely requires a nonlinear learner; the KAN advantage is not an artifact of an easy problem. *Outer controller*. `fedkan_opt` and `kan_full`




Fig. 3. NMSE (mean±std, ten seeds). Both KAN variants are well below the MLPs.

share the identical KAN inner model and differ only by the bandit sparsifier; the move from 9.0 to 4.8 Mb at a small NMSE change isolates the communication-accuracy trade-off to the controller, not the surrogate. *FL heterogeneity*. The proximal term (`mlp_prox` vs. `mlp`) does not close the gap to the KAN, indicating the deficit is the model class rather than the federated-averaging algorithm. *Self-evolution*. `kan_evolve` differs from `fedkan_opt` only by online grid growth; it does not improve accuracy (analyzed next), isolating self-evolution as the inactive factor in this regime.

C. Online behavior under drift

Fig. 5 tracks the global NMSE across slots for a representative seed as interference regimes switch. The KAN surrogates stay consistently below the MLP throughout the non-stationary horizon and recover quickly after a regime change: the federated aggregation re-fits the shifted slice of $p^*(\cdot)$ within a few slots, and the change detector flags the switch with a short delay (reported as detection delay in the released logs). Fig. 4 confirms the trend is not seed-specific: averaged over all seeds, the KAN curves sit below the MLP across the whole horizon, with the proposed sparsified variant tracking the dense KAN closely at half the uplink. Two behaviours are worth noting. First, the error spikes at regime boundaries are shallower and shorter for the KAN, consistent with its smoother inductive bias generalizing across neighbouring regimes rather than overfitting the current one. Second, the curves are stable across seeds (narrow band), so the ranking is a genuine effect and not a lucky run, in line with our multi-seed protocol.

D. Interpretability

Fig. 6 plots the learned univariate responses of `kan_full` to each state variable together with the best-fit closed form (mean fidelity $R^2 = 0.74$), and Table IV reports the dominant form per feature. The optimal power depends near-linearly on the circuit term p_0 ($R^2 = 0.99$) and on interference I ($R^2 = 0.88$), but nonlinearly on channel gain g ($R^2 = 0.36$




Fig. 4. Convergence and tracking under drift: NMSE vs. slot averaged over all seeds (shaded band = $\pm$ std for the proposed method). The KAN advantage is consistent across seeds and across the entire non-stationary horizon, and the proposed sparsified variant tracks the dense KAN at half the uplink.




Fig. 5. Global NMSE over the non-stationary horizon (representative seed).

TABLE IV
DOMINANT EXTRACTED FORM AND SINGLE-BASIS FIDELITY PER INPUT
(FROM KAN_FULL; REFRESHED FROM THE SCALED RUN).

Input	dominant form	R^2
channel gain g	nonlinear	0.36
interference I	near-linear	0.88
circuit power p_0	linear	0.99

for a single basis), which the spline captures. These rules are a direct, auditable output of the model: an operator reads, for instance, that allocated power scales smoothly but sub-linearly with channel gain and decreases with interference, and can sanity-check the policy without instrumenting a black box. An MLP of the same accuracy offers no comparable artifact.

E. Does self-evolution help?

No. `kan_evolve` matches neither `kan_full` (NMSE 0.080 vs. 0.058, $p = 0.014$) nor improves on the proposed static-grid `fedkan_opt`, while using more parameters (404 vs. 360). In this federated streaming regime, where each slot

TABLE V
IN-DOMAIN VS. EXTRAPOLATION NMSE (MEAN OVER TEN SEEDS).
LOWER IS BETTER.

	in-domain	unseen regime (extrap.)
KAN	0.003	0.37
MLP	0.005	2.55

carries limited data, a well-sized static grid already suffices and growing it adds variance rather than capacity that pays off. We report this as a boundary of the “evolutionary capacity” idea.

F. Generalization to unseen interference regimes

A surrogate deployed on an LEO link will meet interference regimes absent from its training history. We train on a subset of regimes ($I \in \{0.2, 1.0\}$) and test on an *unseen* regime ($I = 3.0$). Table V reports in-domain and extrapolation NMSE over ten seeds. In-domain the two surrogates are comparable, but under extrapolation the KAN degrades far more gracefully than the MLP: its additive univariate splines extend smoothly beyond the training support, whereas the MLP’s entangled features extrapolate poorly. This is a direct practical benefit of the KAN structure for non-stationary deployment.

The gap is large (about an order of magnitude in NMSE) and, importantly, the MLP’s extrapolation error has high variance across seeds: occasionally it diverges entirely on the unseen regime, the worst possible behaviour for an autonomous controller that may face a never-seen interference level after a handover. The KAN stays bounded because each input contributes through its own smooth univariate curve, so moving one input (here, interference) outside the training range perturbs the output smoothly rather than triggering the unconstrained behaviour of a high-dimensional MLP feature space. For an LEO deployment, where the set of interference regimes encountered over a mission is not known in advance, this bounded extrapolation is arguably as important as the in-domain accuracy.




Fig. 6. Extracted univariate rules (solid) and best-fit closed form (dashed) for the learned optimal-power policy. The p_0 and I dependences are near-linear; the g dependence is nonlinear and carried by the spline, which an operator can read directly off the trained model.




Fig. 7. Communication-accuracy frontier from a keep-fraction sweep; the bandit controller lands on a favorable knee automatically.

G. Communication-accuracy frontier

Sweeping the update keep-fraction traces the full frontier of the dense KAN (Fig. 7); the accuracy is essentially flat down to a keep-fraction of about 0.5 and then collapses, so roughly half of the uplink is redundant. The bandit-controlled `fedkan_opt` lands in this favorable region automatically. We note plainly that a well-chosen fixed keep-fraction can match or slightly beat the bandit operating point; the controller's value is that it needs no manual tuning and adapts per slot to a time-varying budget $B(t)$, which a fixed setting cannot.

To make the saving concrete, consider the headline operating point: the dense KAN transmits 9.0 Mb over the horizon, whereas the proposed controller transmits 4.8 Mb for the same task, a reduction of roughly one half, at a final NMSE that still undercuts the MLP. Over a constellation with many terminals and short visibility windows this directly relaxes the uplink bottleneck without sacrificing the accuracy advantage. The shape of the frontier also carries an operational lesson: because accuracy is flat above the knee, operators gain almost nothing from spending the full budget, and because it collapses below the knee, over-aggressive compression is dangerous; the learned price keeps the system near the knee as the budget varies.




Fig. 8. (a) accuracy improves with more terminals then saturates; (b) under centralized abundant data a larger grid keeps helping. In the data-limited federated regime, however, growing the grid online overfits (Sec. VI), so a fixed modest grid is preferred.

H. Scaling and the role of capacity

Fig. 8(a) shows accuracy as the number of participating terminals grows: the federated surrogate improves up to about eight to twelve contributors and then saturates, while the uplink grows linearly with participation. Fig. 8(b) sweeps the KAN grid size under *centralized* training with abundant data, where a larger grid monotonically lowers the error, i.e. capacity does help when data is plentiful. This makes the federated result more precise rather than contradicting it: in the federated streaming regime each round carries limited data, so growing the grid online (self-evolution) overfits and does not pay off (Sec. VI), whereas with abundant centralized data extra capacity is useful. The two views together localize *when* capacity helps and explain why a fixed, modest grid is the right choice for the data-limited federated deployment.

I. Per-regime performance and CSI robustness

Two further studies probe robustness. Fig. 9 breaks the final NMSE down by interference regime: the KAN surrogates are at or below the MLP in every regime, so the aggregate advantage is not driven by a single easy operating point but holds across the whole drift range, including the high-interference regime that is hardest to serve. Fig. 10 injects CSI estimation error into the federated loop (the surrogate sees $\hat{g} = g e^\varepsilon$, $\varepsilon \sim \mathcal{N}(0, \sigma_c^2)$, while the oracle uses the true gain) and sweeps σ_c . The KAN retains its edge over the MLP across low-to-moderate CSI error, the operating range of practical channel estimators; under very large error the two converge




Fig. 9. Final NMSE per interference regime. The KAN advantage holds across the drift range, not just on average.




Fig. 10. Robustness to CSI estimation error in the federated loop. The KAN keeps its advantage across low-to-moderate error.

as the input becomes uninformative. This is the deployment-relevant counterpart to the centralized robustness check, and it confirms the advantage is not an artifact of perfect CSI.

VII. DISCUSSION AND LIMITATIONS

The energy-efficiency objective is flat near its optimum, so even imperfect power predictions incur small utility regret; this is why the MLP’s regret is close to the KAN’s despite a much larger prediction error, and why we report both metrics. The interpretability gain is qualitative but concrete: the KAN yields auditable rules, the MLP does not. The KAN advantage is most pronounced in the regime that matters for our setting, namely the data-limited federated/online regime (Sec. VI) and extrapolation to unseen regimes (Table V): the spline inductive bias is sample-efficient and extends smoothly beyond the training support. Given abundant centralized data and in-domain queries an MLP can match the KAN on raw error,

so the benefit is one of sample efficiency, extrapolation, and transparency rather than asymptotic capacity. Our study is deliberately scoped to a single-link, three-variable problem with a simulation oracle so as to isolate and quantify the surrogate’s properties (accuracy, sample efficiency, extrapolation, interpretability) cleanly. The methodology is problem-agnostic, and the most important extensions are to multi-user and multi-sub-band allocation with coupling constraints, comparison against learning-to-optimize baselines such as deep unfolding and WMMSE networks [3], and closing the loop on measured LEO traces. Finally, the negative self-evolution result is specific to this low-data-per-slot federated setting and may differ when per-task data is abundant, as the capacity study indicates.

A. When to prefer the KAN surrogate

The evidence suggests a concrete decision rule for practitioners. Prefer the KAN surrogate when (i) data per update is scarce, as in federated edge deployment with short visibility windows; (ii) the operating distribution can drift into regions not seen during training, where graceful extrapolation matters; or (iii) auditability of the learned policy is required by the operator. When abundant centralized data is available and only in-distribution queries are expected, a plain MLP is an adequate and simpler choice. The communication controller is orthogonal to this choice: it pays off whenever the uplink, not computation, is the binding resource, which is the norm over satellite backhaul.

B. The dual-price view

The controller has a clean economic reading. The multiplier λ_t in (4) is the shadow price of an uplink bit: when the budget $B(t)$ tightens, λ_t rises and the keep-fraction κ_t^* falls, so each terminal transmits only its highest-magnitude (highest-marginal-value) update entries. The bandit estimates this price from realized reward without knowing the drifting distribution or the change points, which is why it tracks the static water-filling rule (5) while remaining robust to non-stationarity. A fixed keep-fraction can match the bandit at a single budget but cannot follow a time-varying $B(t)$, which is the practical regime for shared satellite links.

C. Validity and reproducibility

We guard the conclusions against the failure modes that most often invalidate learned-optimization results. Comparisons are parameter-matched and share the training budget, so the model-class effect is not confounded by capacity or optimization effort; all results are multi-seed with standard deviations and paired significance tests, so no claim rests on a single lucky run; the ablations toggle one component at a time (Table II) so effects are attributable rather than correlational; and we deliberately report the cases where the KAN does *not* win (centralized-abundant accuracy, self-evolution, bandit-vs-best-fixed) rather than omitting them. The oracle is a deterministic optimizer, so labels are exact and the learning target is well-defined. Every figure and table is regenerated from the released result files by a script, and the code reproduces the reported numbers end to end.

VIII. CONCLUSION

We presented an interpretable federated KAN surrogate for autonomous and evolutive resource optimization over non-stationary LEO links. Starting from a hardened mixed-integer formulation (P1), we decomposed the problem into an inner solution-map learning task and an outer budgeted-communication control task, gave the closed-form static rule for the latter, and let an online bandit learn its dual price. Empirically the KAN surrogate learns the energy-efficiency solution map about $2.3\times$ more accurately than a parameter-matched MLP, halves the uplink through learned sparsification while still beating the MLP, extrapolates far more gracefully to unseen interference regimes, and yields closed-form design rules an MLP cannot. We also reported, transparently, that online grid self-evolution does not help in this data-limited federated regime, sharpening the understanding of where evolutive capacity pays off. Future work targets multi-user and multi-sub-band allocation with coupling constraints, comparison against deep unfolding and graph-based learn-to-optimize baselines, deployment-time CSI-error hardening, and validation on measured LEO traces. Interpretable surrogates are, we believe, a practical and trustworthy building block for autonomous optimization in next-generation networked-AI systems.

APPENDIX A

DERIVATION OF THE STATIC KEEP-FRACTION (5)

Consider one slot and drop the slot index. Let terminal n have local update $\Delta\theta_n \in \mathbb{R}^{|\theta|}$ and mask $m_n \in \{0,1\}^{|\theta|}$, with bit cost $c(m_n) = \beta\|m_n\|_0$. Zeroing the entries in the complement of m_n raises the local objective by $\delta_n(m_n)$. To first order around the dense update, the loss increase from dropping entry j is monotone in its magnitude $|\Delta\theta_{n,j}|$ (a coordinate with larger update has larger effect on the model), so at any fixed cardinality $\|m_n\|_0 = k_n$, δ_n is minimized by the top- k_n rule that keeps the k_n largest-magnitude entries. It remains to choose $\{k_n\}$. Relaxing the budget with multiplier $\lambda \geq 0$ gives the separable Lagrangian

$$\mathcal{L} = \sum_n \left[\delta_n(m_n) + \lambda\beta\|m_n\|_0 \right] - \lambda B, \quad (7)$$

whose stationarity in the (sorted) magnitudes keeps entry j iff its marginal loss reduction exceeds the per-bit price, $|\Delta\theta_{n,j}| \geq \tau$ with a common threshold τ set by λ [22]. When the budget binds, the total number of kept entries equals the budget in words, $\sum_n a_n \kappa |\theta| = B/\beta$, and assuming a common keep-fraction κ across the $\sum_n a_n$ participating terminals yields

$$\kappa^* = \min\left(1, \frac{B}{\beta|\theta|\sum_n a_n}\right), \quad (8)$$

which is (5). The proposed controller learns τ (equivalently κ^* , equivalently the price λ) online from the regret-minus-bits reward, avoiding the need to know the drifting distribution or $\Delta\theta$ statistics in advance.

ACKNOWLEDGMENT

This work was supported by the Posts and Telecommunications Institute of Technology (PTIT).

DATA AND CODE AVAILABILITY

The complete source code (problem oracle, KAN/MLP/linear surrogates, federated online loop, bandit controller, drift detector, depth studies, and figure-generation scripts), configuration, and result CSVs that reproduce every number, table, and figure in this paper are openly available at <https://github.com/haodpsut/evokan-leo-opt>.

REFERENCES

- [1] O. Kotheli *et al.*, “Satellite communications in the new space era: A survey and future challenges,” *IEEE Commun. Surveys Tuts.*, vol. 23, no. 1, pp. 70–109, 2021.
- [2] Y. Su, Y. Liu, Y. Zhou, J. Yuan, H. Cao, and J. Shi, “Low earth orbit satellite constellations for global communications,” *IEEE Netw.*, vol. 33, no. 1, pp. 70–77, 2019.
- [3] H. Sun, X. Chen, Q. Shi, M. Hong, X. Fu, and N. D. Sidiropoulos, “Learning to optimize: Training deep neural networks for interference management,” *IEEE Trans. Signal Process.*, vol. 66, no. 20, pp. 5438–5453, 2018.
- [4] B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. A. y Arcas, “Communication-efficient learning of deep networks from decentralized data,” in *AISTATS*, 2017, pp. 1273–1282.
- [5] O. Shahid, S. Pouriyeh, R. M. Parizi, Q. Z. Sheng, G. Srivastava, and L. Zhao, “Communication-efficient federated learning: A survey,” *arXiv preprint arXiv:2107.10996*, 2021.
- [6] Z. Liu, Y. Wang, S. Vaidya, F. Ruehle, J. Halverson, M. Soljačić, T. Y. Hou, and M. Tegmark, “KAN: Kolmogorov-arnold networks,” *arXiv preprint arXiv:2404.19756*, 2024.
- [7] A. Balatsoukas-Stimming and C. Studer, “Deep unfolding for communications systems: A survey and some new directions,” in *IEEE Int. Workshop Signal Process. Syst. (SiPS)*, 2019, pp. 266–271.
- [8] Y. Shen, Y. Shi, J. Zhang, and K. B. Letaief, “Graph neural networks for scalable radio resource management: Architecture design and theoretical analysis,” *IEEE J. Sel. Areas Commun.*, vol. 39, no. 1, pp. 101–115, 2021.
- [9] M. Eisen and A. Ribeiro, “Optimal wireless resource allocation with random edge graph neural networks,” *IEEE Trans. Signal Process.*, vol. 68, pp. 2977–2991, 2020.
- [10] A. Zappone, M. Di Renzo, and M. Debbah, “Wireless networks design in the era of deep learning: Model-based, AI-based, or both?” *IEEE Trans. Commun.*, vol. 67, no. 10, pp. 7331–7376, 2019.
- [11] T. O’Shea and J. Hoydis, “An introduction to deep learning for the physical layer,” *IEEE Trans. Cogn. Commun. Netw.*, vol. 3, no. 4, pp. 563–575, 2017.
- [12] A. N. Kolmogorov, “On the representation of continuous functions of many variables by superposition of continuous functions of one variable and addition,” *Dokl. Akad. Nauk SSSR*, vol. 114, no. 5, pp. 953–956, 1957.
- [13] Blealtan, “efficient-kan: An efficient implementation of kolmogorov-arnold network,” <https://github.com/Blealtan/efficient-kan>, 2024.
- [14] T. Li, A. K. Sahu, M. Zaheer, M. Sanjabi, A. Talwalkar, and V. Smith, “Federated optimization in heterogeneous networks,” *MLSys*, 2020.
- [15] S. Niknam, H. S. Dhillon, and J. H. Reed, “Federated learning for wireless communications: Motivation, opportunities, and challenges,” *IEEE Commun. Mag.*, vol. 58, no. 6, pp. 46–51, 2020.
- [16] M. M. Amiri and D. Gündüz, “Federated learning over wireless fading channels,” *IEEE Trans. Wireless Commun.*, vol. 19, no. 5, pp. 3546–3557, 2020.
- [17] J. Konečný, H. B. McMahan, F. X. Yu, P. Richtárik, A. T. Suresh, and D. Bacon, “Federated learning: Strategies for improving communication efficiency,” *arXiv preprint arXiv:1610.05492*, 2016.
- [18] J. Gama, I. Žliobaitė, A. Bifet, M. Pechenizkiy, and A. Bouchachia, “A survey on concept drift adaptation,” *ACM Comput. Surv.*, vol. 46, no. 4, pp. 1–37, 2014.
- [19] M. De Lange, R. Aljundi, M. Masana, S. Parisot, X. Jia, A. Leonardis, G. Slabaugh, and T. Tuytelaars, “A continual learning survey: Defying forgetting in classification tasks,” *IEEE Trans. Pattern Anal. Mach. Intell.*, vol. 44, no. 7, pp. 3366–3385, 2022.
- [20] C.-X. Wang *et al.*, “Artificial intelligence enabled wireless networking for 5G and beyond: Recent advances and future challenges,” *IEEE Wireless Commun.*, vol. 27, no. 1, pp. 16–23, 2020.

- [21] A. Zappone and E. Jorswieck, “Energy efficiency in wireless networks via fractional programming theory,” *Found. Trends Commun. Inf. Theory*, vol. 11, no. 3–4, pp. 185–396, 2015.
- [22] S. Boyd and L. Vandenberghe, *Convex Optimization*. Cambridge Univ. Press, 2004.
- [23] L. Li, W. Chu, J. Langford, and R. E. Schapire, “A contextual-bandit approach to personalized news article recommendation,” in *WWW*, 2010, pp. 661–670.
- [24] R. S. Sutton and A. G. Barto, *Reinforcement Learning: An Introduction*, 2nd ed. MIT Press, 2018.
- [25] E. S. Page, “Continuous inspection schemes,” *Biometrika*, vol. 41, no. 1/2, pp. 100–115, 1954.

Phuc Hao Do received the Ph.D. degree in telecommunications. He is with the Faculty of Information Technology, Da Nang Architecture University, Da Nang, Vietnam. His research interests include networked AI, federated and quantum-inspired learning, and signal processing for non-terrestrial networks.

Truong Duy Dinh is with the Faculty of Information Security, Posts and Telecommunications Institute of Technology, Ha Noi, Vietnam. His research interests include the Internet of Things, network optimization, and machine learning for communications.